1. Introduction

- By Central Limit Theorem, if sample $(x_1, ..., x_n)$ with each $X_i \sim D(\mu, \sigma^2)$ (D is some distribution) then as $n \to \infty$,

$$\overline{X} \sim N\left(\mu, \frac{\sigma^2}{n}\right)$$

  So distribution of sample mean always tends to normal distribution, with standard deviation $\sigma \,/\, \sqrt{n}$.
- **Unbiased estimate of standard deviation of sample mean**:

$$s = \sqrt{\frac{1}{n-1} \sum_{i=1}^{n} (x_i - \overline{x})^2}$$

- **Standard error of sample mean**: estimate of standard deviation of sample mean: $s \,/\, \sqrt{n}$.
- If n too small then s is poor estimator and mean may not be normally distributed.
- If population distribution is normal and n small then sample mean is t-distributed:

$$\frac{X - \mu}{s \,/\, \sqrt{n}} \sim t_{n-1}$$

  $\frac{X-\mu}{s/\sqrt{n}}$ is **pivotal quantity** as distribution doesn't depend on parameters of X.
- **Hypothesis test** for $\underline{x}$:
  - Define **null hypothesis** which identifies distribution believed to have generated each x_i.
  - Choose **test statistic** h (function of $\underline{x}$), extreme when null is false, not extreme when null is true.
  - **Observed test statistic** is $t = h(\underline{x})$.
  - Determine how extreme t is as a realisation of $T = h(X_1, ..., X_N)$ (so need to know distribution of T).
- **One sided p-value**:

$$\mathbb{P}(T \geq t \mid H_0 \text{ true}) \quad \text{or} \quad \mathbb{P}(T \leq t \mid H_0 \text{ true})$$

- **Two sided p-value**:

$$\mathbb{P}(T \geq |t| \cup T \leq -|t| \mid H_0 \text{ true})$$

2. Monte Carlo testing

- **Monte Carlo testing**: given observed test stat $t = h(\underline{x})$ distribution $F(x \mid \theta)$ hypotheses $H_0 : \theta = \theta_0$, $H_1 : \theta > \theta_0$:
  - For $j \in \{1, ..., N\}$:
    - Simulate n observations $(z_1, ..., z_n)$ from $F(\cdot \mid \theta_0)$.
    - Compute $t_j = h(z_1, ..., z_n)$.
  - Estimate p-value by

$$P(T \geq t \mid H_0 \text{ true}) \approx \frac{1}{N} \sum_{j=1}^{N} \mathbb{I}\{t_j \geq t\}$$

3. The bootstrap

- **The non-parametric bootstrap estimate**: given independent data $\underline{x} = (x_1, ..., x_n)$ and stat $S(\cdot)$, **resample** (draw samples of size n with replacement) $\underline{x}$ B times to give $\underline{x}^{*1}, ..., \underline{x}^{*B}$. To compute **bootstrap estimate of standard error of S**, compute

$$\widehat{\operatorname{Var}}(S(\underline{x})) = \frac{1}{B-1} \sum_{b=1}^{B} \left(S\left(\underline{x}^{*b}\right) - \overline{S}^*\right)^2$$

  where

$$\overline{S}^* = \frac{1}{B} \sum_{b=1}^{B} S\left(\underline{x}^{*b}\right)$$

  The standard error estimate is then $\sqrt{\widehat{\operatorname{Var}}(S(\underline{x}))}$, i.e. the standard deviation of $S\left(\underline{x}^{*1}\right), ..., S\left(\underline{x}^{*B}\right)$ The **bootstrap estimate** of S is simply $S(\underline{x})$.
- For random variable X, **(cumulative) distribution function (cdf)** $F : \mathbb{R} \to [0, 1]$ is

$$F_X(x) = F(x) := \mathbb{P}(X \leq x)$$

- Properties of cdf:
  - $\lim_{x \to -\infty} F(x) = 0$ and $\lim_{x \to \infty} F(x) = 1$.
  - **Monotonicity**: $x' < x \Longrightarrow F(x') \leq F(x)$.
  - **Right-continuity**: $\lim_{t \to x^+} F(t) = F(x)$.
- Given data $(x_1, ..., x_n)$ with each sample i.i.d. realisation of random variable X, **empirical (cumulative) distribution function (ecdf)** is

$$\hat{F}(x) := \frac{1}{n} \sum_{i=1}^{n} \mathbb{I}\{x_i \leq x\}$$

- Let $X_1, ..., X_n$ be random sample from distribution with cdf F. Then

$$\sup_{x \in \mathbb{R}} \left|\hat{F}(x) - F(x)\right| \to 0 \quad \text{as } n \to \infty$$

- Given data $(x_1, ..., x_n)$, sampling uniformly at random from $\underline{x}$ is equivalent to sampling from distribution with cdf defined as ecdf constructed from $\underline{x}$.
- For mean of sample of m draws from ecdf constructed from n data points, expectation and variance are

$$\mathbb{E}\left[\overline{Y}\right] = \overline{x}, \quad \operatorname{Var}\left(\overline{Y}\right) = \frac{n-1}{n} \frac{s_x^2}{m}$$

- So $\widehat{\operatorname{Var}}(S(\underline{x}) \to \frac{n-1}{n} \frac{s^2}{n}$ as $B \to \infty$.
- If **sampling fraction** $f = \frac{n}{N}$ where N population size, n sample size, is $f \geq 0.1$, can't assume infinite population.
- Given finite population of size N, mean $\overline{X}$ of sample drawn uniformly at random without replacement has variance

$$\mathrm{Var}(\overline{X}) = \frac{N-n}{N-1}\frac{\sigma^2}{n}$$

where σ^2 is true population variance.

- Given finite population of size N, sample of size n with variance S^2 drawn without replacement,

$$\mathbb{E}\left[\left(1 - \frac{n}{N}\right)\frac{S^2}{n}\right] = \mathrm{Var}(\overline{X})$$

  so it is unbiased estimator of $\mathrm{Var}(\overline{X})$
- **Population bootstrap**: given independent data $(x_1, ..., x_n)$ drawn from finite population of size N, assuming $N / n = k$ is integer, construct new data set

$$\underline{\tilde{x}} = (x_1, ..., x_n, x_1, ..., x_n, ..., x_1, ..., x_n)$$

  by repeating $\underline{x}$ k times. Then construct B new samples $\underline{x}^{*1}, ..., \underline{x}^{*B}$ by sampling without replacement. Then compute

$$\widehat{\mathrm{Var}}(S(\underline{x})) = \frac{1}{B-1}\sum_{b=1}^{B}\left(S(\underline{x}^{*b}) - \overline{S}^*\right)^2$$

  where

$$\overline{S}^* = \frac{1}{B}\sum_{b=1}^{B} S(\underline{x}^{*b})$$

  If N / n not integer, $N = kn + m$ for $0 < m < n$, then before each of the B samples, append to $\underline{\tilde{x}}$ a sample without replacement of size m from $\underline{x}$.
- If data believed to follow type of distribution, can use **parametric bootstrap**: given independent data $(x_1, ..., x_n)$, believed to be drawn from distribution $F(\cdot, \theta)$ with parameter θ:
  - Find maximum likelihood estimator $\hat{\theta}$.
  - Draw B new samples of size n from $F\left(\cdot, \hat{\theta}\right)$ to give $\underline{x}^{*1}, ..., \underline{x}^{*B}$.
  - Compute

$$\widehat{\mathrm{Var}}(S(\underline{x})) = \frac{1}{B-1}\sum_{b=1}^{B}\left(S(\underline{x}^{*b}) - \overline{S}^*\right)^2$$

    where

$$\overline{S}^* = \frac{1}{B}\sum_{b=1}^{B} S(\underline{x}^{*b})$$

- For parameter θ of distribution, estimated by statistic S, with $\hat{\theta} = S(\underline{x})$, **bias** is

$$\mathrm{bias}\left(\theta, \hat{\theta}\right) = \mathbb{E}\left[\hat{\theta}\right] - \theta$$

- **Basic bootstrap bias estimate**:

$$\widehat{\operatorname{bias}}\left(\theta, \hat{\theta}\right) = \overline{S}^{*} - \hat{\theta} = \frac{1}{B}\sum_{b=1}^{B} S\left(\underline{x}^{*b}\right) - S(\underline{x})$$

- **Bias correction**: subtract bias from usual estimate:

$$\hat{\theta} - \widehat{\operatorname{bias}}\left(\theta, \hat{\theta}\right) = 2\hat{\theta} - \overline{S}^{*}$$

  But often $2\hat{\theta} - \overline{S}^{*}$ has higher variance as estimator than $\hat{\theta}$.
- **Normal confidence interval for bootstrap estimate**:

$$\hat{\theta} \pm z_{\alpha/2}\sqrt{\widehat{\operatorname{Var}}(S(\underline{x}))}$$

  where $z_{\alpha/2}$ is $100(\alpha / 2)\%$ percentile of standard normal distribution. **Note**: only valid if size of data large enough, need to check for normality of bootstrap samples using quantile plot.
- **Percentile confidence interval**: use if $\hat{F}$ close to true distribution. $100(1-\alpha)\%$ confidence interval is

$$\left[S^{*}_{((\alpha/2)B)}, S^{*}_{((1-\alpha/2)B)}\right]$$

  where $S^{*}_{(i)}$ is ith largest value of $S\left(\underline{x}^{*b}\right)$ for $b = 1, ..., B$. B must be chosen to make $(\alpha / 2)B$ and $(1 - \alpha / 2)B$ integers. B must be > 2000 for this to be good estimate. **Note**: inaccurate if bias or non-constant standard error or distribution of $S(X) \mid \theta$ isn't symmetric.
- **BC (bias corrected)** and **BCa (bias corrected and accelerated)** confidence intervals make adjustments when bias is present or there is non-constant standard error.

4. Monte Carlo integration

- Let random variable Y take values in sample space Ω with pdf f_Y, then

$$\mu := \mathbb{E}[Y] = \int_{\Omega} y f_Y(y) \, \mathrm{d}y$$

- μ approximated by

$$\hat{\mu}_n = \frac{1}{n}\sum_{i=1}^{n} Y_i$$

  for i.i.d. samples Y_i.
- If $Y = g(X)$ with X random variable with pdf f_X, then

$$\mu = \mathbb{E}[Y] = \mathbb{E}[g(X)] = \int g(x) f_X(x) \, \mathrm{d}x$$

- To estimate $\int_a^b f(x) \, \mathrm{d}x$, use $X \sim$ Unif(a, b)

$$\mu = \int_a^b f(x) \, \mathrm{d}x = \int_a^b (b-a) f(x) \frac{1}{b-a} = \int_a^b (b-a) f(x) f_X(x) = \mathbb{E}[(b-a) f(X)]$$

  which can be estimated by

$$\hat{\mu}_n = (b-a)\frac{1}{n}\sum_{i=1}^{n} g(X_i)$$

for i.i.d. samples X_i.

- If $\mathrm{Var}(Y) = \sigma^2 < \infty$, Monte Carlo integration unbiased as $\mathbb{E}\left[\hat{\mu}_n\right] = \mu$.
- **Mean-square error**: $\mathrm{Var}\left(\hat{\mu}_n\right) = \mathbb{E}\left[\left(\hat{\mu}_n - \mu\right)^2\right] = \frac{\sigma^2}{n}$.
- **Root mean-square error**: $\mathrm{RMSE} = \sqrt{\mathbb{E}\left[\left(\hat{\mu}_n - \mu\right)^2\right]} = \frac{\sigma}{\sqrt{n}}$.
- RMSE is $O\left(n^{-1/2}\right)$.
- For functions f, g, $f(n) = O(g(n))$ as $n \to \infty$ if exist $C, n_0 \in \mathbb{R}$ such that

$$\forall n \geq n_0, \quad |f(n)| \leq Cg(n)$$

- **Midpoint Riemann integral estimate**:

$$\int_a^b f(x)\,\mathrm{d}x = \frac{b-a}{n}\sum_{i=1}^{n} f(x_i)$$

where

$$x_i = a + \frac{b-a}{n}\left(i - \frac{1}{2}\right)$$

- For d dimensions, Riemann sum converges in $O\left(n^{-2/d}\right)$, Monte Carlo converges in $O\left(n^{-1/2}\right)$ regardless of d.
- $100(1-\alpha)\%$ confidence interval for Monte Carlo integration:

$$\mu \in \hat{\mu}_n \pm z_{\alpha/2}\frac{\sigma}{\sqrt{n}}$$

where σ estimated with standard sample deviation of $\left\{y_i\right\} = \{g(x_i)\}$.

- If $g(x)$ constant multiple of indicator function, $g(x) = c\mathbb{I}\{A(x)\}$ for condition A, then

$$\hat{p}_n = \frac{1}{n}\sum_{i=1}^{n}\mathbb{I}\{A(x_i)\}$$

is estimator for $p = \mathbb{P}(A)$. Binomial confidence interval is

$$p \in \hat{p}_n \pm z_{\alpha/2}\sqrt{\frac{\hat{p}_n\left(1-\hat{p}_n\right)}{n}}$$

- TODO: fill in notes here
- Probability of no 1s in n Monte Carlo samples is $(1-p)^n$ so one-sided $100(1-\alpha)\%$ confidence interval has upper bound $p \leq 1 - a^{1/n} \approx -\frac{\log(a)}{n}$ using Taylor expansion.
- If $\hat{p}$ very small and non-zero,

$$cz_{\alpha/2}\sqrt{\frac{\hat{p}_n\left(1-\hat{p}_n\right)}{n}} \approx cz_{\alpha/2}\sqrt{\frac{\hat{p}_n}{n}}$$

so relative error is

$$\delta := cz_{\alpha/2}\sqrt{\frac{\hat{p}_n}{n}}\ /\ \hat{p} = \frac{cz_{\alpha/2}}{\sqrt{\hat{p}_n n}}$$

for relative error at most δ,

$$n \geq \frac{c^2 z_{\alpha/2}^2}{\hat{p}_n \delta^2}$$

so n grows inversely with $\hat{p}_n$.
- To estimate probability of event $\mathbb{P}(X \in E)$, Monte Carlo estimate $\mathbb{E}[\mathbb{I}\{X \in E\}]$.

5. Simulation

- Let F cdf, then **generalised inverse cdf** is

$$F^{-1}(u) := \inf\{x : F(x) \geq u\}$$

- **Inverse transform sampling algorithm**: let random variable X with cdf F, with generalised inverse F^{-1}.
  - Simulate $U \sim \mathrm{Unif}(0, 1)$.
  - Compute $X \sim F^{-1}(U)$.

  X is then distributed with cdf F. Only works for 1D distributions.
- **Rejection sampling algorithm**: given **target density** function f, **proposal density** function $\tilde{f}$ with $\forall x \in \mathbb{R}^d$, $f(x) \leq c\tilde{f}(x)$ for some $c < \infty$,
  - Set $a = \text{false}$
  - While $a = \text{false}$:
    - Simulate $u \sim \mathrm{Unif}(0, 1)$.
    - Simulate $x \sim \tilde{f}(\cdot)$.
    - If $u \leq \frac{f(x)}{c\tilde{f}(x)}$, set $a = \text{true}$.
  - Once while loop exited, return x, which is distributed with pdf f.
- **Note**: f and $\tilde{f}$ don't need to be normalised.
- When $f, \tilde{f}$ normalised, expected number of iterations of rejection sampling algorithm is c.
- **Important**: when choosing value of c, always round **up** if inexact.
- When checking if rejection sampling can be used, check if ratio $f(x)\ /\ \tilde{f}(x)$ tends to 0 as $x \to \pm\infty$ and differentiate ratio with respect to x to find maximum.